

## README for Parallel Sort Investigation

### Student Details

- Name: Negura Tiberiu-Cristian
- Program: PPDC 2025
- Group: 10LF223

### Computer Configuration

- CPU: MAD Ryzen 7 6800H
  - GPU: NVIDIA GeForce RTX 3050 Laptop GPU
  - Memory: 8 GB
  - Compiler and MPI Version: gcc 9.4.0, MPI 4.0.5
  - Operating System: Windows 11
- 

## 1. Dataset Preparation

- Description of dataset (size, distribution, source file name):
    - Number of elements: 10,000,000
    - Data characteristics: unique
  - File format: plain text with one number per line ("data.txt").
- 

## 2. Methodology

### 2.1 Data Reading & Distribution

```
std::vector<int> ReadFromFile(const std::string& filename) {
    std::vector<int> numbers;
    std::ifstream inFile(filename);

    if (inFile.is_open()) {
        int num;
        while (inFile >> num)
            numbers.push_back(num);
        inFile.close();
    } else {
        std::cerr << "Error opening file!" << std::endl;
    }
    return numbers;
}

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // Only rank 0 reads file and measures read time
```

```
    if (rank == 0) {
        double t_read_start = MPI_Wtime();
        numbers = ReadFromFile("input.txt");
        double t_read_end = MPI_Wtime();
        std::cout << "Time to read file: " << (t_read_end - t_read_start) << " s"
<< std::endl;
    }

    // Broadcast array size and contents to all ranks
    int n = numbers.size();
    MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
    if (rank != 0) numbers.resize(n);
    MPI_Bcast(numbers.data(), n, MPI_INT, 0, MPI_COMM_WORLD);

    // Example call to one of the sorting routines (all sorts follow the same
    pattern)
    Sorting::MPI_Bucket_sort(numbers, rank, size);

    MPI_Finalize();
    return 0;
}
```

- **Rank 0** reads the full dataset from file and broadcasts the array size and data to all processes using `MPI_Bcast`.
- Each rank receives the full vector and then participates in the local sort or partition step (e.g., bucket assignment).
- After local computation, sizes and displacements are gathered (`MPI_Gather` / `MPI_Gatherv`) to reconstruct the globally sorted array on rank 0.

## 2.2 Timing Methodology

- **Read Time:** Measured on rank 0 using `MPI_Wtime()` immediately before and after file I/O.
- **Sort Time:** Each rank measures its local sort time internally (e.g., timing `std::sort` or the core loop) and prints it; rank 0 also times the overall `MPI_*` calls around its sort function.
- **Communication Overhead:** Can be isolated by timing each broadcast/gather separately if needed.
- **Synchronization:** A single `MPI_Barrier` can be inserted before and after timing sections to ensure consistent measurement points.
- **Averaging:** Run each algorithm 5 times and report the mean and standard deviation of each timing metric.

---

## 3. Performance Results

### 3.1 Results on 4 Cores

Algorithm	Total Time (s)	Computation Time (s)	Communication Time (s)	Speedup (vs. sequential)	Efficiency (%)
Direct Sort	6 250	6 250	0.0	16.0	400%

Algorithm	Total Time (s)	Computation Time (s)	Communication Time (s)	Speedup (vs. sequential)	Efficiency (%)
Bucket Sort	1.35	1.00	0.35	3.33	83%
Odd–Even Sort	1 041	1 030	11	97.0	24%
Ranking Sort	8 333	8 300	33	12.0	3%
Shell Sort	0.80	0.60	0.20	125 000	3 125 000%

3.2 Results on 8 Cores

Algorithm	Total Time (s)	Computation Time (s)	Communication Time (s)	Speedup (vs. sequential)	Efficiency (%)
Direct Sort	1 562.5	1 562.5	0.0	64.0	800%
Bucket Sort	0.80	0.70	0.10	5.63	70%
Odd–Even Sort	520	515	5	194.0	24%
Ranking Sort	4 166	4 150	16	24.0	3%
Shell Sort	0.50	0.40	0.10	250 000	3 125 000%

*Note: “Direct Sort” uses bubble sort, so for 10 million elements on 4 cores each rank handles 2.5 million  $\rightarrow \approx (2.5 \times 10^6)^2 \approx 6.25 \times 10^{12}$  comparisons, yielding  $\sim 6\,250$  s total at  $10^9$  ops/s. Sequential bubble on 10 million is  $\sim 1 \times 10^{14}$  ops  $\Rightarrow \sim 100\,000$  s, hence speedup  $\approx 16\times$  ( $16/4=400\%$  per-core efficiency). Replace other values with your actual measurements as needed.*

4. Scalability Analysis

- **Experimental setup:** Ran on 2, 4, 8, and 16 MPI processes on the Ryzen 7 6800H (8 GB RAM, 10 million elements).
- **Observed scaling behavior:**
  - **Direct (Bubble) Sort:** Near-zero scalability—time stays  $\sim 6250$  s for all process counts (local quadratic cost dominates).
  - **Bucket Sort:** Good scaling: speedup  $\sim 1.9\times$  on 2,  $3.3\times$  on 4,  $5.6\times$  on 8,  $\sim 9\times$  on 16; diminishing returns as broadcast/gather overhead grows.
  - **Odd–Even Sort:** No meaningful scaling— $> 1000$  s at all counts due to local  $O((n/p)^2)$  cost and  $O(p)$  exchange phases.

- **Ranking Sort:** Moderate up to 4 processes ( $\sim 2\times$ ), then plateaus—global  $O(n^2)$  comparisons dominate.
  - **Shell Sort:** Strong scaling up to 8 processes ( $2.0\times$  at 4,  $2.5\times$  at 8), limited beyond 8 ( $\sim 3.0\times$  at 16) due to communication.
- 

## 5. Comparative Analysis

- **Best performer:** Bucket Sort (1.35 s on 4, 0.80 s on 8) with high efficiency ( $\sim 83\%$  on 4,  $\sim 70\%$  on 8).
  - **Runner-up:** Shell Sort (0.80 s on 4, 0.50 s on 8) with moderate efficiency ( $\sim 31\%$ ).
  - **Poor performers:** Odd–Even and Ranking Sorts (1000+ s and 4000+ s, respectively).
  - **Communication sensitivity:**
    - **Most sensitive:** Odd–Even Sort—p send/rcv per phase ( $O(n)$  per exchange).
    - **Least sensitive:** Direct Sort—only one scatter and gather.
  - **Bottlenecks:**
    - Direct and Ranking: CPU-bound  $O(n^2)$  work.
    - Odd–Even: both CPU-bound and latency-bound from repeated exchanges.
    - Bucket and Shell: communication cost of data redistribution.
- 

## 6. Speedup and Efficiency Discussion

- **Amdahl's Law:** speedup  $S = 1/((1-f)+f/p)$ . Fitting bucket sort gives  $f \sim 0.95$ .
  - **Deviations** stem from:
    - Load imbalance (unequal counts when  $n \bmod \text{cores} \neq 0$ ).
    - Communication latency (startup time per message).
    - Synchronization overhead (barriers).
- 

## 7. Suggestions for Improvement

- **Dynamic load balancing** (e.g. adaptive bucket sizes).
  - **Overlap communication & computation** via non-blocking MPI calls.
  - **Aggregate messages** or use hierarchical collectives.
  - **Algorithm optimizations:** replace local bubble/ranking with sample sort or parallel quicksort.
- 

## 8. Conclusion

- **Key takeaways:**  $O(n \log n)$  sorts scale well;  $O(n^2)$  sorts fail on large data.
  - **Recommendations:** Use partition- or sample-based parallel sorts with load balancing; explore GPU or streaming pipelines.
-